

# Quantifying the World

## Case Study 2

Presented by:

Alexander Sepenu

Nnenna Okpara

Taifur Chowdhury

Juan Edgar Nunez - Gonzalez

## Business Understanding

# Quantifying the World

Objective: The objective of this case study is to build a classification model using logistic regression which predicts hospital readmittance.

```
In [1]: import os
os.chdir(r'/Users/juannunez/Documents/SMU/Quantifying The World/Case study 2')
```

```
In [2]: import pandas as pd
import numpy as np
import pandas_profiling
from pandas_profiling import ProfileReport
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

```
In [3]: # reading in the data
dai_df = pd.read_csv('diabetic_data.csv')
```

In [4]:

```
# visualizing the shape of the data for rows and column  
print("Diabetes Data Shape is: {}".format(dai_df.shape))
```

Diabetes Data Shape is: (101766, 50)

# Data Evaluation / Engineering

## Exploratory Data Analysis

### Summary

The data contains missing values which are represented with a question mark. Hence, we looked for where and how many values are missing in the dataset. To facilitate the analysis of this data we will replace the "?" with a NaN.

Next we identified columns with diagnostic codes ( i.e. "diag\_1", "diag\_2", "diag\_3" ) that are mostly numeric, with some exceptions where the code is alphanumeric. Since, we don't have a way to identify the meaning of the diagnostic codes, we elected to remove the alphanumeric anomalies (i.e., "V10"), etc.; Because these values can affect the entire Logistic Regression, if not corrected.

In the third step, we computed the percentages of missing values in every column. Our team considered to drop the columns where more than 20% of the values were missing, and impute the missing values in columns where less than 10% of the values were missing. However, with since there is no absolute rule of thumb, we settled on a more conservative approach based on James Ledoux <https://jamesrledoux.com/code/imputation>, experience and dropped any column with more than 5% of values missing and imputed columns with less than 5% of values missing.

We use the Mode to impute the missing values because it is the value less likely to skew the data.

In [5]:

```
# Looking at the first column of data to see the distribution of data across  
dai_df.head()
```

Out [5]:

	encounter_id	patient_nbr	race	gender	age	weight	admission_type_id	discharge
0	2278392	8222157	Caucasian	Female	[0-10)	?		6
1	149190	55629189	Caucasian	Female	[10-20)	?		1
2	64410	86047875	AfricanAmerican	Female	[20-30)	?		1
3	500364	82442376	Caucasian	Male	[30-40)	?		1
4	16680	42519267	Caucasian	Male	[40-50)	?		1

5 rows x 50 columns

In [6]:

```
# changing "?" character to NAN to have an over view of missing values within
dai_df = dai_df.replace('?', np.nan)
```

In [7]:

```
# looking at the dataset data type
dai_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101766 entries, 0 to 101765
Data columns (total 50 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   encounter_id                          101766 non-null int64
1   patient_nbr                           101766 non-null int64
2   race                                  99493 non-null  object
3   gender                                101766 non-null object
4   age                                   101766 non-null object
5   weight                                3197 non-null   object
6   admission_type_id                     101766 non-null int64
7   discharge_disposition_id              101766 non-null int64
8   admission_source_id                   101766 non-null int64
9   time_in_hospital                      101766 non-null int64
10  payer_code                             61510 non-null  object
11  medical_specialty                     51817 non-null  object
12  num_lab_procedures                    101766 non-null int64
13  num_procedures                        101766 non-null int64
14  num_medications                       101766 non-null int64
15  number_outpatient                      101766 non-null int64
16  number_emergency                       101766 non-null int64
17  number_inpatient                      101766 non-null int64
```

```

18  diag_1                101745 non-null object
19  diag_2                101408 non-null object
20  diag_3                100343 non-null object
21  number_diagnoses      101766 non-null int64
22  max_glu_serum         101766 non-null object
23  A1Cresult             101766 non-null object
24  metformin             101766 non-null object
25  repaglinide           101766 non-null object
26  nateglinide           101766 non-null object
27  chlorpropamide        101766 non-null object
28  glimepiride           101766 non-null object
29  acetohexamide         101766 non-null object
30  glipizide             101766 non-null object
31  glyburide             101766 non-null object
32  tolbutamide           101766 non-null object
33  pioglitazone          101766 non-null object
34  rosiglitazone         101766 non-null object
35  acarbose              101766 non-null object
36  miglitol              101766 non-null object
37  troglitazone          101766 non-null object
38  tolazamide            101766 non-null object
39  examide               101766 non-null object
40  citoglipton           101766 non-null object
41  insulin               101766 non-null object
42  glyburide-metformin   101766 non-null object
43  glipizide-metformin   101766 non-null object
44  glimepiride-pioglitazone 101766 non-null object
45  metformin-rosiglitazone 101766 non-null object
46  metformin-pioglitazone 101766 non-null object
47  change                101766 non-null object
48  diabetesMed           101766 non-null object
49  readmitted            101766 non-null object
dtypes: int64(13), object(37)
memory usage: 38.8+ MB

```

```

In [8]: # Data set head with NaN to see the distribution of missing values in data
dai_df.head()

```

Out [8]:

	encounter_id	patient_nbr	race	gender	age	weight	admission_type_id	discharge
0	2278392	8222157	Caucasian	Female	[0-10)	NaN		6
1	149190	55629189	Caucasian	Female	[10-20)	NaN		1
2	64410	86047875	AfricanAmerican	Female	[20-30)	NaN		1
3	500364	82442376	Caucasian	Male	[30-40)	NaN		1
4	16680	42519267	Caucasian	Male	[40-50)	NaN		1

5 rows x 50 columns

In [9]:

```
# Fishing out the number of missing values per column as part of the EDA  
dai_df.isnull().sum()
```

```

Out[9]: encounter_id      0
        patient_nbr      0
        race              2273
        gender            0
        age               0
        weight            98569
        admission_type_id 0
        discharge_disposition_id 0
        admission_source_id 0
        time_in_hospital  0
        payer_code        40256
        medical_specialty  49949
        num_lab_procedures 0
        num_procedures     0
        num_medications    0
        number_outpatient  0
        number_emergency   0
        number_inpatient   0
        diag_1             21
        diag_2             358
        diag_3             1423
        number_diagnoses   0
        max_glu_serum      0
        Alcresult          0
        metformin          0
        repaglinide        0
        nateglinide        0
        chlorpropamide     0
        glimepiride        0
        acetohexamide      0
        glipizide          0
        glyburide          0
        tolbutamide        0
        pioglitazone       0
        rosiglitazone      0
        acarbose           0
        miglitol           0
        troglitazone       0
        tolazamide         0
        examide            0
        citoglipton        0
        insulin            0
        glyburide-metformin 0
        glipizide-metformin 0
        glimepiride-pioglitazone 0
        metformin-rosiglitazone 0
        metformin-pioglitazone 0
        change             0
        diabetesMed        0
        readmitted         0
        dtype: int64

```

```
In [10]: # Using regex function to correct wrongly entered values for diag_1, diag_2,

dai_df["diag_1"] =dai_df["diag_1"].replace(to_replace ='[V]', value = '', reg
dai_df["diag_1"] = dai_df["diag_1"].replace(to_replace ='[E]', value = '', re

dai_df["diag_2"] =dai_df["diag_2"].replace(to_replace ='[V]', value = '', reg
dai_df["diag_2"] = dai_df["diag_2"].replace(to_replace ='[E]', value = '', re

dai_df["diag_3"] =dai_df["diag_3"].replace(to_replace ='[V]', value = '', reg
dai_df["diag_3"] = dai_df["diag_3"].replace(to_replace ='[E]', value = '', re
```

```
In [11]: # Calculating the percentages of missing values per column to direct our impu
((dai_df.isnull() | dai_df.isna()).sum() * 100 / dai_df.index.size).round(2)
```

```
Out[11]: encounter_id      0.00
patient_nbr      0.00
race             2.23
gender           0.00
age              0.00
weight           96.86
admission_type_id 0.00
discharge_disposition_id 0.00
admission_source_id 0.00
time_in_hospital 0.00
payer_code       39.56
medical_specialty 49.08
num_lab_procedures 0.00
num_procedures   0.00
num_medications  0.00
number_outpatient 0.00
number_emergency 0.00
number_inpatient 0.00
diag_1           0.02
diag_2           0.35
diag_3           1.40
number_diagnoses 0.00
max_glu_serum    0.00
AlCresult        0.00
metformin        0.00
repaglinide      0.00
nateglinide      0.00
chlorpropamide   0.00
glimepiride      0.00
acetohexamide    0.00
glipizide        0.00
glyburide        0.00
tolbutamide      0.00
pioglitazone     0.00
rosiglitazone    0.00
acarbose         0.00
miglitol         0.00
troglitazone     0.00
tolazamide       0.00
examide          0.00
citoglipton      0.00
insulin          0.00
glyburide-metformin 0.00
glipizide-metformin 0.00
glimepiride-pioglitazone 0.00
metformin-rosiglitazone 0.00
metformin-pioglitazone 0.00
change           0.00
diabetesMed      0.00
readmitted       0.00
dtype: float64
```



```
In [12]: #dropping unwanted columns with large missing values  
new_diab = dai_df.drop(['weight', 'payer_code', 'medical_specialty' ], axis = 1)
```

```
In [13]: # imputing nan vallues  
new_diab = new_diab.apply(lambda x: x.fillna(x.value_counts().index[0]))
```

```
In [14]: #Confirming that there are no missing values  
((new_diab.isnull() | new_diab.isna()).sum() * 100 / new_diab.index.size).rou
```

```

Out[14]: encounter_id      0.0
         patient_nbr      0.0
         race             0.0
         gender           0.0
         age              0.0
         admission_type_id 0.0
         discharge_disposition_id 0.0
         admission_source_id 0.0
         time_in_hospital  0.0
         num_lab_procedures 0.0
         num_procedures    0.0
         num_medications   0.0
         number_outpatient  0.0
         number_emergency   0.0
         number_inpatient  0.0
         diag_1            0.0
         diag_2            0.0
         diag_3            0.0
         number_diagnoses  0.0
         max_glu_serum     0.0
         A1Cresult         0.0
         metformin         0.0
         repaglinide       0.0
         nateglinide       0.0
         chlorpropamide    0.0
         glimepiride       0.0
         acetohexamide     0.0
         glipizide         0.0
         glyburide         0.0
         tolbutamide       0.0
         pioglitazone      0.0
         rosiglitazone     0.0
         acarbose          0.0
         miglitol          0.0
         troglitazone      0.0
         tolazamide        0.0
         examide           0.0
         citoglipton       0.0
         insulin           0.0
         glyburide-metformin 0.0
         glipizide-metformin 0.0
         glimepiride-pioglitazone 0.0
         metformin-rosiglitazone 0.0
         metformin-pioglitazone 0.0
         change            0.0
         diabetesMed       0.0
         readmitted        0.0
         dtype: float64

```

```

In [15]: # Visualizing basic statistics for the data
         new_diab.describe()

```

Out[15]:

	encounter_id	patient_nbr	admission_type_id	discharge_disposition_id	admission_s
<b>count</b>	1.017660e+05	1.017660e+05	101766.000000	101766.000000	101766
<b>mean</b>	1.652016e+08	5.433040e+07	2.024006	3.715642	5
<b>std</b>	1.026403e+08	3.869636e+07	1.445403	5.280166	4
<b>min</b>	1.252200e+04	1.350000e+02	1.000000	1.000000	1
<b>25%</b>	8.496119e+07	2.341322e+07	1.000000	1.000000	1
<b>50%</b>	1.523890e+08	4.550514e+07	1.000000	1.000000	1
<b>75%</b>	2.302709e+08	8.754595e+07	3.000000	4.000000	1
<b>max</b>	4.438672e+08	1.895026e+08	8.000000	28.000000	25

## Detail EDA with vizualization

We are recoding all categorical variables into numeric to do a thorough EDA and feature analysis.

In [16]:

```
# convert categorical to integers
cleanup_nums = {"race": {"AfricanAmerican": 0, "Asian": 1, "Caucasian": 2, "H
"gender": {"Female": 0, "Male": 1, "Unknown/Invalid": 2},
"age": {"[0-10)": 0, "[10-20)": 1, "[20-30)": 2, "[30-40)": 3, "[
"max_glu_serum": {">200": 0, ">300": 1, ">300": 2, "None": 3,
"AlCresult": {">7": 0, ">8": 1, "None": 2, "Norm": 4},
"metformin": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"repaglinide": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"nateglinide": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"chlorpropamide": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"glimepiride": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"acetohexamide": {"No": 0, "Steady": 1},
"glipizide": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"glyburide": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"tolbutamide": {"No": 0, "Steady": 1},
"pioglitazone": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"rosiglitazone": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"acarbose": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"miglitol": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"troglitazone": {"No": 0, "Steady": 1},
"tolazamide": {"No": 0, "Steady": 1, "Up": 2},
"examide": {"No": 0},
"citoglipton": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"insulin": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"glyburide-metformin": {"Down": 0, "No": 1, "Steady": 2, "Up": 3},
"glipizide-metformin": {"No": 0, "Steady": 1},
"glimepiride-pioglitazone": {"Down": 0, "No": 1, "Steady": 2, "U
"metformin-rosiglitazone": {"No": 0, "Steady": 1},
"metformin-pioglitazone": {"No": 0, "Steady": 1},
"change": {"No": 0, "Ch": 1},
"diabetesMed": {"No": 0, "Yes": 1}}
```

In [17]:

```
#Integrating all recorded numeric features into dataframes
new_diab = new_diab.replace(cleanup_nums)
```

In [18]:

```
#Changing float data type into intergers
new_diab['diag_1'] = new_diab['diag_1'].astype(float).astype(int)
new_diab['diag_2'] = new_diab['diag_2'].astype(float).astype(int)
new_diab['diag_3'] = new_diab['diag_3'].astype(float).astype(int)
new_diab['gender'] = new_diab['gender'].astype(int)
```

In [19]:

```
target_labels = new_diab['readmitted'].unique().tolist()
```

```
In [20]: target_mod = {  
          'NO':0,  
          '>30':0,  
          '<30':1  
        }
```

```
In [21]: new_diab['readmitted_binary'] = new_diab['readmitted'].map(target_mod)  
new_diab['readmitted_binary'].unique()
```

```
Out[21]: array([0, 1])
```

```
In [22]: #dropping unwanted columns  
new_diab = new_diab.drop(['readmitted'], axis = 1)  
new_diab = new_diab.drop(['examide'], axis = 1)  
new_diab = new_diab.drop(['citoglipton'], axis = 1)
```

```
In [24]: #Dropping target feature for feature analysis  
new_diab_df = new_diab.drop(['readmitted_binary'], axis = 1)
```

```
In [25]: new_diab_df.dtypes
```

```
Out[25]: encounter_id      int64
patient_nbr      int64
race             int64
gender           int64
age             int64
admission_type_id int64
discharge_disposition_id int64
admission_source_id int64
time_in_hospital int64
num_lab_procedures int64
num_procedures    int64
num_medications   int64
number_outpatient int64
number_emergency  int64
number_inpatient  int64
diag_1            int64
diag_2            int64
diag_3            int64
number_diagnoses  int64
max_glu_serum     int64
A1Cresult         int64
metformin         int64
repaglinide       int64
nateglinide       int64
chlorpropamide    int64
glimepiride       int64
acetoexamide      int64
glipizide         int64
glyburide         int64
tolbutamide       int64
pioglitazone      int64
rosiglitazone     int64
acarbose          int64
miglitol          int64
troglitazone      int64
tolazamide        int64
insulin           int64
glyburide-metformin int64
glipizide-metformin int64
glimepiride-pioglitazone int64
metformin-rosiglitazone int64
metformin-pioglitazone int64
change            int64
diabetesMed        int64
dtype: object
```

```
In [ ]:  #(new_diab_df.profile_report()).to_file('Diabetes.html')
```

```
In [27]:  #diab_profile = ProfileReport(new_diab, title="Profiling Report", explorative
```

```
In [29]: #diab_profile.to_widgets()
```

## EDA Output

EDA Output from Profile Package Detailing the Outline of the data

### Overview

Overview Alerts 53 Reproduction	
Dataset statistics	
Number of variables	47
Number of observations	101766
Missing cells	0
Missing cells (%)	0.0%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	36.5 MiB
Average record size in memory	376.0 B
Variable types	
Numeric	16
Categorical	28
Boolean	3

The above image details the outline of the data showing the various data types in the data. The data has 101766 rows and 50 columns of originally but 3 columns were dropped due significant missing values being more that our adpated appraoch of dropping 5% or more missing values. There are 47 columns left for further analysis of which 28 are categorical, 3 being boolean (true or false) and 16 are numerical.

## EDA Output Cont

EDA Output from the Profile Package Detailing the Outline of the variables in the data

## Variables

encounter\_id

Real number ( $\mathbb{R}_{>0}$ )

HIGH CORRELATION

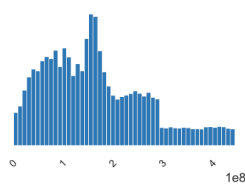
HIGH CORRELATION

HIGH CORRELATION

UNIQUE

Distinct	101766
Distinct (%)	100.0%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	165201645.6

Minimum	12522
Maximum	443867222
Zeros	0
Zeros (%)	0.0%
Negative	0
Negative (%)	0.0%
Memory size	795.2 KiB



1e8

Toggle details

patient\_nbr

Real number ( $\mathbb{R}_{>0}$ )

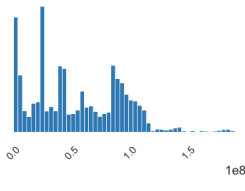
HIGH CORRELATION

HIGH CORRELATION

HIGH CORRELATION

Distinct	71518
Distinct (%)	70.3%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	54330400.69

Minimum	135
Maximum	189502619
Zeros	0
Zeros (%)	0.0%
Negative	0
Negative (%)	0.0%
Memory size	795.2 KiB



1e8

Toggle details

race

Categorical

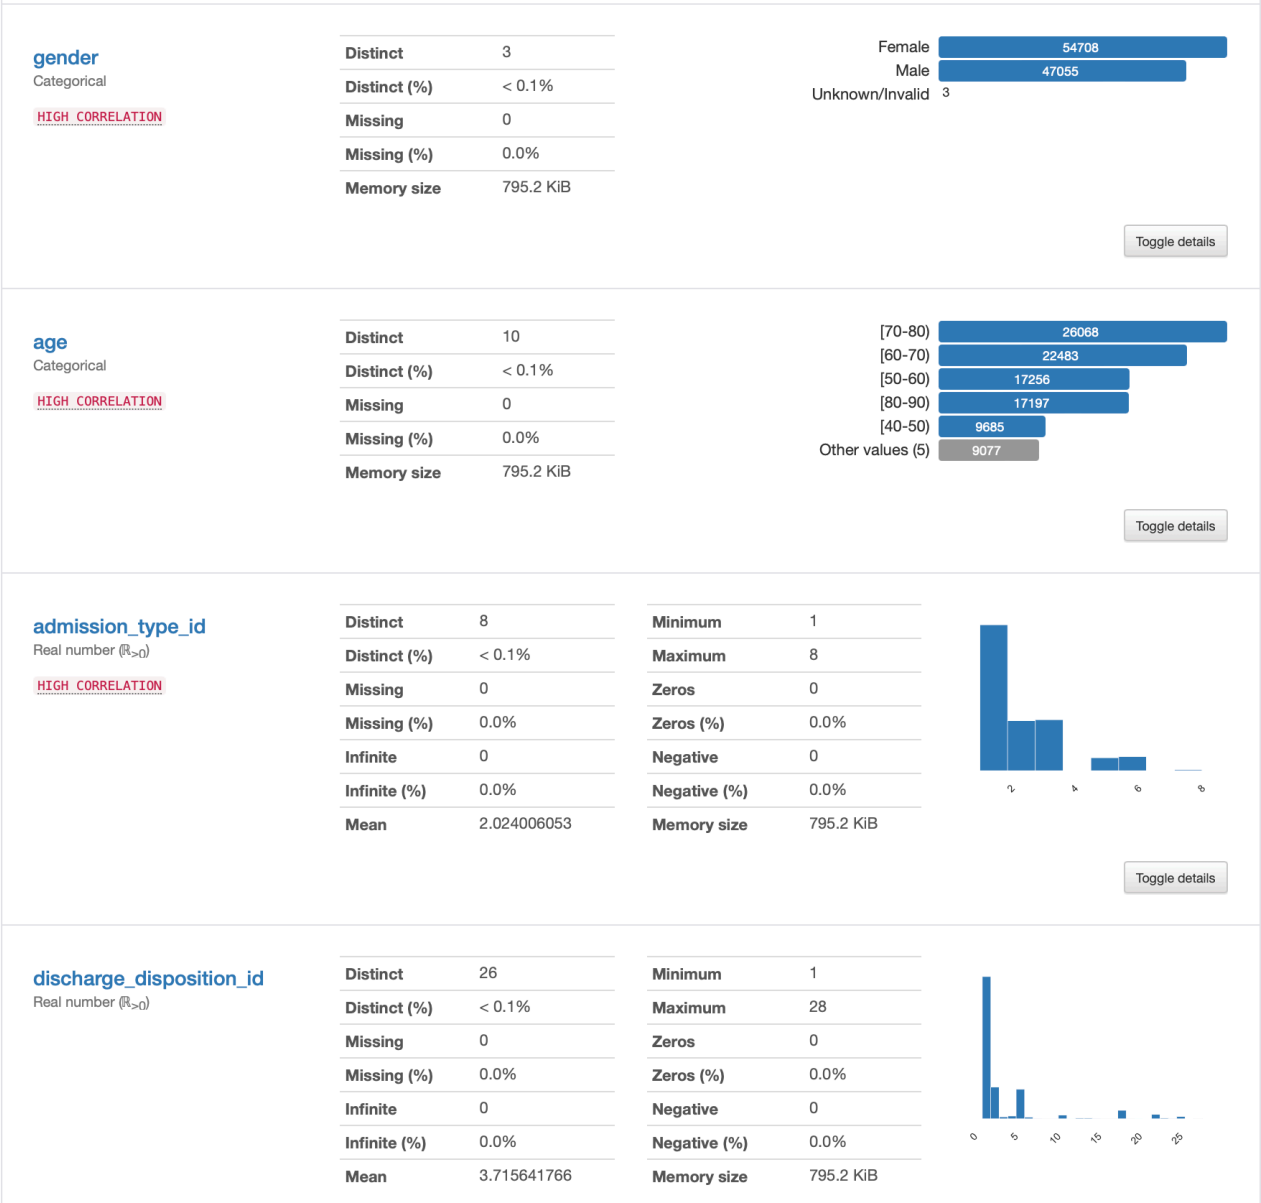
HIGH CORRELATION

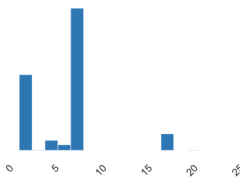
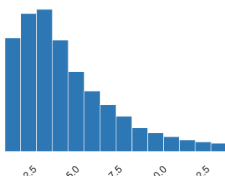
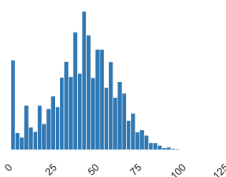
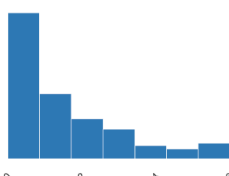
Distinct	5
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	795.2 KiB

Caucasian	78372
AfricanAmerican	19210
Hispanic	2037
Other	1506
Asian	641

Toggle details





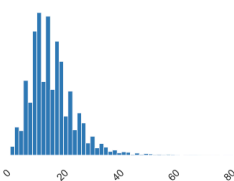
<div>admission_source_id</div> <div>Real number (<math>\mathbb{R}_{&gt;0}</math>)</div> <div>HIGH CORRELATION</div>	<table><tr><td>Distinct</td><td>17</td></tr><tr><td>Distinct (%)</td><td>&lt; 0.1%</td></tr><tr><td>Missing</td><td>0</td></tr><tr><td>Missing (%)</td><td>0.0%</td></tr><tr><td>Infinite</td><td>0</td></tr><tr><td>Infinite (%)</td><td>0.0%</td></tr><tr><td>Mean</td><td>5.754436649</td></tr></table>	Distinct	17	Distinct (%)	< 0.1%	Missing	0	Missing (%)	0.0%	Infinite	0	Infinite (%)	0.0%	Mean	5.754436649	<table><tr><td>Minimum</td><td>1</td></tr><tr><td>Maximum</td><td>25</td></tr><tr><td>Zeros</td><td>0</td></tr><tr><td>Zeros (%)</td><td>0.0%</td></tr><tr><td>Negative</td><td>0</td></tr><tr><td>Negative (%)</td><td>0.0%</td></tr><tr><td>Memory size</td><td>795.2 KiB</td></tr></table>	Minimum	1	Maximum	25	Zeros	0	Zeros (%)	0.0%	Negative	0	Negative (%)	0.0%	Memory size	795.2 KiB	 <div>Toggle details</div>
Distinct	17																														
Distinct (%)	< 0.1%																														
Missing	0																														
Missing (%)	0.0%																														
Infinite	0																														
Infinite (%)	0.0%																														
Mean	5.754436649																														
Minimum	1																														
Maximum	25																														
Zeros	0																														
Zeros (%)	0.0%																														
Negative	0																														
Negative (%)	0.0%																														
Memory size	795.2 KiB																														
<div>time_in_hospital</div> <div>Real number (<math>\mathbb{R}_{&gt;0}</math>)</div>	<table><tr><td>Distinct</td><td>14</td></tr><tr><td>Distinct (%)</td><td>&lt; 0.1%</td></tr><tr><td>Missing</td><td>0</td></tr><tr><td>Missing (%)</td><td>0.0%</td></tr><tr><td>Infinite</td><td>0</td></tr><tr><td>Infinite (%)</td><td>0.0%</td></tr><tr><td>Mean</td><td>4.395986872</td></tr></table>	Distinct	14	Distinct (%)	< 0.1%	Missing	0	Missing (%)	0.0%	Infinite	0	Infinite (%)	0.0%	Mean	4.395986872	<table><tr><td>Minimum</td><td>1</td></tr><tr><td>Maximum</td><td>14</td></tr><tr><td>Zeros</td><td>0</td></tr><tr><td>Zeros (%)</td><td>0.0%</td></tr><tr><td>Negative</td><td>0</td></tr><tr><td>Negative (%)</td><td>0.0%</td></tr><tr><td>Memory size</td><td>795.2 KiB</td></tr></table>	Minimum	1	Maximum	14	Zeros	0	Zeros (%)	0.0%	Negative	0	Negative (%)	0.0%	Memory size	795.2 KiB	 <div>Toggle details</div>
Distinct	14																														
Distinct (%)	< 0.1%																														
Missing	0																														
Missing (%)	0.0%																														
Infinite	0																														
Infinite (%)	0.0%																														
Mean	4.395986872																														
Minimum	1																														
Maximum	14																														
Zeros	0																														
Zeros (%)	0.0%																														
Negative	0																														
Negative (%)	0.0%																														
Memory size	795.2 KiB																														
<div>num_lab_procedures</div> <div>Real number (<math>\mathbb{R}_{&gt;0}</math>)</div>	<table><tr><td>Distinct</td><td>118</td></tr><tr><td>Distinct (%)</td><td>0.1%</td></tr><tr><td>Missing</td><td>0</td></tr><tr><td>Missing (%)</td><td>0.0%</td></tr><tr><td>Infinite</td><td>0</td></tr><tr><td>Infinite (%)</td><td>0.0%</td></tr><tr><td>Mean</td><td>43.09564098</td></tr></table>	Distinct	118	Distinct (%)	0.1%	Missing	0	Missing (%)	0.0%	Infinite	0	Infinite (%)	0.0%	Mean	43.09564098	<table><tr><td>Minimum</td><td>1</td></tr><tr><td>Maximum</td><td>132</td></tr><tr><td>Zeros</td><td>0</td></tr><tr><td>Zeros (%)</td><td>0.0%</td></tr><tr><td>Negative</td><td>0</td></tr><tr><td>Negative (%)</td><td>0.0%</td></tr><tr><td>Memory size</td><td>795.2 KiB</td></tr></table>	Minimum	1	Maximum	132	Zeros	0	Zeros (%)	0.0%	Negative	0	Negative (%)	0.0%	Memory size	795.2 KiB	 <div>Toggle details</div>
Distinct	118																														
Distinct (%)	0.1%																														
Missing	0																														
Missing (%)	0.0%																														
Infinite	0																														
Infinite (%)	0.0%																														
Mean	43.09564098																														
Minimum	1																														
Maximum	132																														
Zeros	0																														
Zeros (%)	0.0%																														
Negative	0																														
Negative (%)	0.0%																														
Memory size	795.2 KiB																														
<div>num_procedures</div> <div>Real number (<math>\mathbb{R}_{&gt;0}</math>)</div> <div>ZEROS</div>	<table><tr><td>Distinct</td><td>7</td></tr><tr><td>Distinct (%)</td><td>&lt; 0.1%</td></tr><tr><td>Missing</td><td>0</td></tr><tr><td>Missing (%)</td><td>0.0%</td></tr><tr><td>Infinite</td><td>0</td></tr><tr><td>Infinite (%)</td><td>0.0%</td></tr><tr><td>Mean</td><td>1.339730362</td></tr></table>	Distinct	7	Distinct (%)	< 0.1%	Missing	0	Missing (%)	0.0%	Infinite	0	Infinite (%)	0.0%	Mean	1.339730362	<table><tr><td>Minimum</td><td>0</td></tr><tr><td>Maximum</td><td>6</td></tr><tr><td>Zeros</td><td>46652</td></tr><tr><td>Zeros (%)</td><td>45.8%</td></tr><tr><td>Negative</td><td>0</td></tr><tr><td>Negative (%)</td><td>0.0%</td></tr><tr><td>Memory size</td><td>795.2 KiB</td></tr></table>	Minimum	0	Maximum	6	Zeros	46652	Zeros (%)	45.8%	Negative	0	Negative (%)	0.0%	Memory size	795.2 KiB	 <div>Toggle details</div>
Distinct	7																														
Distinct (%)	< 0.1%																														
Missing	0																														
Missing (%)	0.0%																														
Infinite	0																														
Infinite (%)	0.0%																														
Mean	1.339730362																														
Minimum	0																														
Maximum	6																														
Zeros	46652																														
Zeros (%)	45.8%																														
Negative	0																														
Negative (%)	0.0%																														
Memory size	795.2 KiB																														

num\_medications

Real number ( $\mathbb{R}_{\geq 0}$ )

Distinct	75
Distinct (%)	0.1%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	16.02184423

Minimum	1
Maximum	81
Zeros	0
Zeros (%)	0.0%
Negative	0
Negative (%)	0.0%
Memory size	795.2 KIB



Toggle details

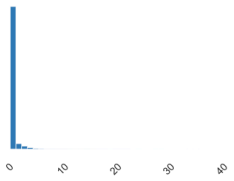
number\_outpatient

Real number ( $\mathbb{R}_{\geq 0}$ )

ZEROS

Distinct	39
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	0.3693571527

Minimum	0
Maximum	42
Zeros	85027
Zeros (%)	83.6%
Negative	0
Negative (%)	0.0%
Memory size	795.2 KIB



Toggle details

number\_emergency


Real number ( $\mathbb{R}_{\geq 0}$ )

SKEWED

ZEROS

Distinct	33
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	0.1978362125

Minimum	0
Maximum	76
Zeros	90383
Zeros (%)	88.8%
Negative	0
Negative (%)	0.0%
Memory size	795.2 KIB



Toggle details

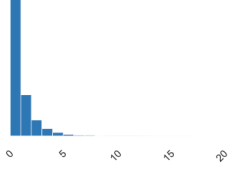
number\_inpatient

Real number ( $\mathbb{R}_{\geq 0}$ )

ZEROS

Distinct	21
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	0.6355659061

Minimum	0
Maximum	21
Zeros	67630
Zeros (%)	66.5%
Negative	0
Negative (%)	0.0%
Memory size	795.2 KIB



Toggle details

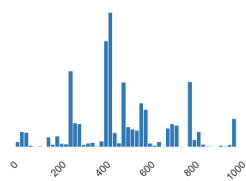
diag\_1

Real number ( $\mathbb{R}_{>0}$ )

HIGH CORRELATION

Distinct	709
Distinct (%)	0.7%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	486.5183199

Minimum	3
Maximum	999
Zeros	0
Zeros (%)	0.0%
Negative	0
Negative (%)	0.0%
Memory size	795.2 KiB



Toggle details

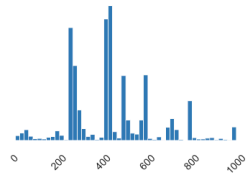
diag\_2

Real number ( $\mathbb{R}_{>0}$ )

diag\_2HIGH CORRELATION

Distinct	705
Distinct (%)	0.7%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	434.4104617

Minimum	2
Maximum	999
Zeros	0
Zeros (%)	0.0%
Negative	0
Negative (%)	0.0%
Memory size	795.2 KiB



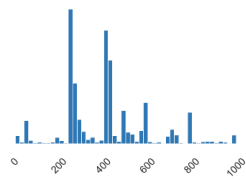
Toggle details

diag\_3

Real number ( $\mathbb{R}_{>0}$ )

Distinct	726
Distinct (%)	0.7%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	407.4365624

Minimum	1
Maximum	999
Zeros	0
Zeros (%)	0.0%
Negative	0
Negative (%)	0.0%
Memory size	795.2 KiB



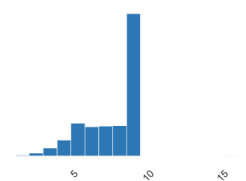
Toggle details

number\_diagnoses

Real number ( $\mathbb{R}_{>0}$ )

Distinct	16
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	7.422606765

Minimum	1
Maximum	16
Zeros	0
Zeros (%)	0.0%
Negative	0
Negative (%)	0.0%
Memory size	795.2 KiB



Toggle details

metformin-pioglitazone

Categorical

HIGH CORRELATION

Distinct	2
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	795.2 KiB

No101765

Steady1

Toggle details

change

Categorical

HIGH CORRELATION

HIGH CORRELATION

Distinct	2
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	795.2 KiB

No54755

Ch47011

Toggle details

diabetesMed

Boolean

HIGH CORRELATION

HIGH CORRELATION

Distinct	2
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	99.5 KiB

True78363

False23403

Toggle details

readmitted

Categorical

HIGH CORRELATION

Distinct	3
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	795.2 KiB

NO54864

>3035545

<3011357

Toggle details


The outline shows that the data is heavily imbalanced (skewed data). These imply that data needs to be scaled by standardization by centering the data and dividing by the standard deviation to shift the distribution to have a mean of zero and a standard deviation of one for model interpretability.

Source: <https://machinelearningmastery.com/standardscaler-and-minmaxscaler-transforms-in-python/>

## Feature Analysis using Correlation Matrix

Output from the Profile Package Detailing the Correlation Variables with Significance

# Correlations



The above correlation matrix show some of the strong correlations between different features. As expected patient ID and patient number are correlated. time\_spent\_in\_hospital is positively correlated to patients\_time in\_lab and number\_of\_medication . Also, number\_of\_medication is positively correlated to number\_of\_procedure. number\_of\_emergency\_visits is positively correlated to number\_of\_inpatient.

The Profile package was used to select features of relevance based on their correlation, using the Pearson's Correlation Metrics. We have gathered a list of important features for our model building using 21 out of 47 features but this will be probed further to firm up the decision to move forward with this assertion.

Although correlation matrix above can be a useful tool to find multicollinearity, their outcome only shows a bivariate relationship between the independent variables in our dataset, hence we are exploring other mean to check and deal with the problem of multicollinearity. A simple method to detect multicollinearity in a model is to use a Variance Inflation Factor(VIF) approach to get a better understanding at this.

The table below will be our reference point to tackle this problem. From the table, we will drop any VIF above 5 and to be more conservative use all VIF below 2.5 score.

VIF Threshold	Reference Type	Reference Date	Reference
<b>VIF &gt; 10</b> is problematic	Book	2012	Vittinghoff E, Glidden DV, Shiboski SC, McCulloch CE. Regression Methods in Biostatistics: Linear, Logistic, Survival, and Repeated Measures Models. 2nd ed. 2012 edition. Springer; 2011.
<b>VIF &gt; 5</b> or <b>VIF &gt; 10</b> is problematic	Book	2017	James G, Witten D, Hastie T, Tibshirani R. An Introduction to Statistical Learning: With Applications in R. 1st ed. 2013, Corr. 7th printing 2017 edition. Springer; 2013.
<b>VIF &gt; 5</b> is cause for concern and <b>VIF &gt; 10</b> indicates a serious collinearity problem	Book	2001	Menard S. Applied Logistic Regression Analysis. 2nd edition. SAGE Publications, Inc; 2001.
<b>VIF ≥ 2.5</b> indicates considerable collinearity	Research Paper	2018	Johnston R, Jones K, Manley D. Confounding and collinearity in regression analysis: a cautionary tale and an alternative procedure, illustrated by studies of British voting behaviour. Qual Quant. 2018;52(4):1957-1976. doi:10.1007/s11135-017-0584-6

Source:<https://quantifyinghealth.com/vif-threshold/>

## Feature Analysis Checking for Multicollinearity using VIF

Based on the VIF output in checking or detecting multicollinearity and dealing with its effect on our model, we will drop predictors with high collinearity from the dataset based on the desired threshold of at most 2.5 score for modeling.

```
In [30]: # Calculation multicollinearity for each predictor  
X = new_diab
```

```
In [31]: # Import library and setting parameters for VIF computation  
from statsmodels.stats.outliers_influence import variance_inflation_factor  
  
def calculate_vif_(X, thresh=2.5):  
    variables = list(range(X.shape[1]))  
    dropped = True  
    while dropped:  
        dropped = False  
        vif = [variance_inflation_factor(X.iloc[:, variables].values, ix)  
               for ix in range(X.iloc[:, variables].shape[1])]  
  
        maxloc = vif.index(max(vif))  
        if max(vif) > thresh:  
            print('dropping \'' + X.iloc[:, variables].columns[maxloc] +  
                  '\\' at index: ' + str(maxloc))  
            del variables[maxloc]  
            dropped = True  
  
    print('Remaining variables:')  
    print(X.columns[variables])  
    print(X.iloc[:, variables].shape[1])
```

```
In [32]: calculate_vif_(X, thresh=2.5)
```



```

dropping 'glimepiride-pioglitazone' at index: 39
dropping 'miglitol' at index: 33
dropping 'chlorpropamide' at index: 24
dropping 'acarbose' at index: 31
dropping 'glyburide-metformin' at index: 34
dropping 'nateglinide' at index: 23
dropping 'repaglinide' at index: 22
dropping 'max_glu_serum' at index: 19
dropping 'number_diagnoses' at index: 18
dropping 'glimepiride' at index: 20
dropping 'rosiglitazone' at index: 25
dropping 'age' at index: 4
dropping 'pioglitazone' at index: 23
dropping 'glyburide' at index: 21
dropping 'glipizide' at index: 20
dropping 'metformin' at index: 18
dropping 'num_medications' at index: 10
dropping 'AlCresult' at index: 16
dropping 'num_lab_procedures' at index: 8
dropping 'diabetesMed' at index: 24
dropping 'diag_2' at index: 13
dropping 'diag_1' at index: 12
dropping 'diag_3' at index: 12
dropping 'race' at index: 2
dropping 'insulin' at index: 15
dropping 'time_in_hospital' at index: 6
dropping 'admission_type_id' at index: 3
Remaining variables:
Index(['encounter_id', 'patient_nbr', 'gender', 'discharge_disposition_id',
      'admission_source_id', 'num_procedures', 'number_outpatient',
      'number_emergency', 'number_inpatient', 'acetoexamide', 'tolbutamide',
      'troglitazone', 'tolazamide', 'glipizide-metformin',
      'metformin-rosiglitazone', 'metformin-pioglitazone', 'change',
      'readmitted_binary'],
      dtype='object')
18

```

```

In [33]: #Index of usable features
usable_ind_vars = ['encounter_id', 'patient_nbr', 'gender', 'discharge_disposi
      'admission_source_id', 'num_procedures', 'number_outpatient',
      'number_emergency', 'number_inpatient', 'acetoexamide', 'tolbutamide'
      'troglitazone', 'tolazamide', 'glipizide-metformin',
      'metformin-rosiglitazone', 'metformin-pioglitazone', 'change',
      'readmitted_binary']

```

```

In [34]: #dropping unwanted features based on their VIFs criteria defined above
model_diab_df = new_diab.drop(['glimepiride-pioglitazone', 'miglitol', 'chlor
      'repaglinide', 'max_glu_serum', 'number_diagnoses', 'gl
      'glipizide', 'metformin', 'num_medications', 'AlCresult
      'race', 'insulin', 'time_in_hospital', 'admission_type_

```

After VIF computations, these features were dropped based on their high collinearity with other predictors in the dataset. These features are: 'glimepiride-pioglitazone', 'miglitol', 'chlorpropamide', 'acarbose', 'glyburide-metformin', 'nateglinide', 'repaglinide', 'max\_glu\_serum', 'number\_diagnoses', 'glimepiride', 'rosiglitazone', 'age', 'pioglitazone', 'glipizide', 'metformin', 'num\_medications', 'A1Cresult', 'num\_lab\_procedures', 'diabetesMed', 'diag\_2', 'race', 'insulin', 'time\_in\_hospital' and 'admission\_type\_id'.

This gives us a refined dataset to work with in building our model of which a total of 17 out of 49 features from the original dataset will be used after data cleaning and preparation while the target data remains 'readmitted\_binary'. The features are 'encounter\_id', 'patient\_nbr', 'gender', 'discharge\_disposition\_id', 'admission\_source\_id', 'num\_procedures', 'number\_outpatient', 'number\_emergency', 'number\_inpatient', 'acetoheaxamide', 'tolbutamide', 'troglitazone', 'tolazamide', 'glipizide-metformin', 'metformin-rosiglitazone', 'metformin-pioglitazone' and 'change'.

In [35]:

```
# Checking out our refined data set after data cleaning and preparation
model_diab_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101766 entries, 0 to 101765
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   encounter_id                          101766 non-null  int64
1   patient_nbr                           101766 non-null  int64
2   gender                                101766 non-null  int64
3   discharge_disposition_id              101766 non-null  int64
4   admission_source_id                   101766 non-null  int64
5   num_procedures                        101766 non-null  int64
6   number_outpatient                     101766 non-null  int64
7   number_emergency                      101766 non-null  int64
8   number_inpatient                      101766 non-null  int64
9   acetoheaxamide                       101766 non-null  int64
10  tolbutamide                           101766 non-null  int64
11  troglitazone                          101766 non-null  int64
12  tolazamide                            101766 non-null  int64
13  glipizide-metformin                   101766 non-null  int64
14  metformin-rosiglitazone                101766 non-null  int64
15  metformin-pioglitazone                 101766 non-null  int64
16  change                                101766 non-null  int64
17  readmitted_binary                     101766 non-null  int64
dtypes: int64(18)
memory usage: 14.0 MB
```

In [36]:

```
#saving new dataset for modelling
model_diab_df.to_csv('diab_model_df.csv', index = False)
```

# Model Preparation

We intend to evaluate our logistic regression model using three methods - a 70/30 split with no cross validation, K-Fold, and ShuffleSplit techniques.

Our base model will be a 70/30 train-test split to measure the performance of accuracy, F1-score, Precision, Recall, and AUC metrics, and then compare the metrics against the K-Fold and the ShuffleSplit technique.

Given our objective of predicting hospital re-admittance, the chosen method will help us build an effective model. The best competing model will help us to predict our target(re-admittance) response using unseen data.

We will measure the output of each approach to see which one gives us the best value in terms of Accuracy, F1-score, Precision, Recall, AUC metrics to decide which model is better.

To begin this analysis, the dataset needs to be standardized to make sure the content and the format are internally consistent. We standardize data when features have wide differences between ranges. For example, when there are numerical data with different measures (such as weight, distance, etc). the process helps the model to internalize the data and train itself effectively.

Source: <https://builtin.com/data-science/when-and-why-standardize-your-data>

In [37]:

```
# reading in the cleaned data for modeling
model_df = pd.read_csv('diab_model_df.csv')
```

In [38]:

```
model_df.head()
```

Out[38]:

	encounter_id	patient_nbr	gender	discharge_disposition_id	admission_source_id	num_pro
0	2278392	8222157	0	25	1	
1	149190	55629189	0	1	7	
2	64410	86047875	0	1	7	
3	500364	82442376	1	1	7	
4	16680	42519267	1	1	7	

# Model Building & Evaluation

## Model Evaluation

```
In [40]: # First approach is to scale or standardized the data to center the predictor
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

sc = StandardScaler()
X = pd.DataFrame(sc.fit_transform(model_df.drop(['readmitted_binary'], axis =

y = model_df['readmitted_binary']

X.head()
```

```
Out[40]:
```

	0	1	2	3	4	5	6	7	
0	-1.587330	-1.191545	-0.927397	4.031022	-1.169873	-0.785398	-0.291461	-0.21262	-0.50
1	-1.608075	0.033564	-0.927397	-0.514312	0.306482	-0.785398	-0.291461	-0.21262	-0.50
2	-1.608901	0.819654	-0.927397	-0.514312	0.306482	2.145781	1.286748	-0.21262	0.28
3	-1.604653	0.726480	1.078031	-0.514312	0.306482	-0.199162	-0.291461	-0.21262	-0.50
4	-1.609366	-0.305227	1.078031	-0.514312	0.306482	-0.785398	-0.291461	-0.21262	-0.50

```
In [186... # 70/30 train and test split on df
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, ran
```

```
In [183... # 70/30 model with no cross validation using Logistics Regression

classify = LogisticRegression(random_state=0).fit(X_train, y_train)
pred_result = classify.predict_proba(X)
pred_result_class = classify.predict(X)

score = classify.score(X_test, y_test)

print('Accuracy score without validation is : ', score)
```

```
Accuracy score without validation is : 0.8849328529315428
[-2.43819579 -1.95179934 -2.16800043 ... -2.44017292 -2.04895162
-2.48000782]
```

In [178...

```

from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
from sklearn import datasets, metrics, model_selection

predictions = LogisticRegression(random_state=123456).fit(X_test,y_test)
predict_result = predictions.predict_proba(X_test)
predict_result_class = predictions.predict(X_test)

#use model to predict probability that given y value is 1
y_pred = predictions.predict_proba(X_test)[::,1]

#calculate AUC of model
auc = metrics.roc_auc_score(y_test, y_pred)

#print AUC score
print('This is AUC:', auc)

print(classification_report(y_test,predict_result_class))

```

This is AUC: 0.6342195527283687

	precision	recall	f1-score	support
0	0.89	1.00	0.94	27035
1	0.48	0.02	0.03	3495
accuracy			0.89	30530
macro avg	0.69	0.51	0.49	30530
weighted avg	0.84	0.89	0.84	30530

In [179...

```
#Implementing cross validation (cross_val_score) using Logistics Regression
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import confusion_matrix
from sklearn import datasets, linear_model
from sklearn.model_selection import cross_val_score
import sklearn.metrics as metrics

model = LogisticRegression(random_state = 123456).fit(X_train, y_train)

y_pred = cross_val_predict(model, X_train, y_train, cv=10)

scores = cross_val_score(model, X_train, y_train, cv=5)

f1_score = cross_val_score(model, X_train, y_train, cv=5,scoring='f1_macro')
precision = cross_val_score(model, X_train, y_train, cv=5,scoring='precision')
recall = cross_val_score(model, X_train, y_train, cv=5,scoring='recall')
roc_auc = cross_val_score(model, X_train, y_train, cv=5,scoring= 'roc_auc')

print('This is Accuracy:',scores.mean())

print('This is F1_Score:',f1_score.mean())

print('This is Precision:',precision.mean())

print('This is Recall:',recall.mean())

print('This is AUC:',roc_auc.mean())
```

```
This is Accuracy: 0.8894659905941957
This is F1_Score: 0.48584354923615536
This is Precision: 0.479160310588882
This is Recall: 0.01564456824692772
This is AUC: 0.6290506234805731
```

In [194...

```

# Model validation using using ShuffleSplit
#from sklearn.model_selection import ShuffleSplit
#from sklearn.linear_model import LogisticRegression
#from sklearn.metrics import accuracy_score
#X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=1)
#num_cv_iterations = 10
#num_instances = len(y)
#cv_object = ShuffleSplit(n_splits=num_cv_iterations,
#                          test_size = 0.3)
#acc_score_shuf = []

#model = LogisticRegression(random_state = 1)
#for iter_num, (train_indices, test_indices) in enumerate(cv_object.split(X,y)
#    model.fit(X[train_indices],y[train_indices]) # train object
#    y_hat = model.predict(X[test_indices]) # get test set precitions
#    pred_values = model.predict(X_test)
#    acc = accuracy_score(pred_values , y_test)
#    acc_score_shuf.append(acc)

#avg_acc_shuf_score = sum(acc_score_shuf)/num_cv_iterations

#print('accuracy of each fold - {}'.format(acc_score_shuf))
#print('Avg accuracy : {}'.format(avg_acc_shuf_score))

```

```

In [53]: # Model validation using using ShuffleSplit
from sklearn.model_selection import ShuffleSplit

from sklearn.metrics import accuracy_score
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
sum_cv_iterations = 10
sum_instances = len(y)
cv_object = ShuffleSplit(n_splits=num_cv_iterations,
                        test_size = 0.3)
acc_score_shuf = []

model = LogisticRegression(random_state = 0)
for iter_num, (train_indices, test_indices) in enumerate(cv_object.split(X,y)):
    model.fit(X[train_indices],y[train_indices]) # train object
    y_hat = model.predict(X[test_indices]) # get test set precitions
    pred_values = model.predict(X_test)
    acc = accuracy_score(pred_values , y_test)
    acc_score_shuf.append(acc)

avg_acc_shuf_score = sum(acc_score_shuf)/num_cv_iterations

print('accuracy of each fold - {}'.format(acc_score_shuf))
print('Avg accuracy : {}'.format(avg_acc_shuf_score))

accuracy of each fold = [0.8880198097633833, 0.8880198097633833, 0.8880198097633833, 0.8880198097633833, 0.8880198097633833, 0.8880198097633833, 0.8880198097633833, 0.8880198097633833, 0.8880198097633833, 0.8880198097633833]
Avg accuracy : 0.8880198097633833

```

NB: when running this code, it produces an error on this work computer but runs on another computer.

In [105...

```
#from sklearn.neighbors import KNeighborsRegressor

#knn_model = KNeighborsRegressor(n_neighbors=3)

#knn_model.fit(X_train, y_train)

from sklearn.neighbors import KNeighborsClassifier
import sklearn.metrics as metrics
# Create KNN classifier
model = LogisticRegression(random_state = 2)
knn = KNeighborsClassifier(n_neighbors = 8)
# Fit the classifier to the data
knn.fit(X_train,y_train)

y_pred = knn.predict(X_test)

print(metrics.accuracy_score(y_test, y_pred))
print(metrics.classification_report(y_test, y_pred))
```

0.8841794955781199

	precision	recall	f1-score	support
0	0.89	1.00	0.94	27035
1	0.38	0.02	0.03	3495
accuracy			0.88	30530
macro avg	0.63	0.51	0.49	30530
weighted avg	0.83	0.88	0.83	30530

In [134...

```
from sklearn.model_selection import cross_validate
from sklearn.metrics import make_scorer
from sklearn.metrics import confusion_matrix

model = LogisticRegression(random_state = 123456)
def confusion_matrix_scorer(model, X, y):
    y_pred = model.predict(X_train)
    cm = confusion_matrix(y_train, y_test)
    return {'tn': cm[0, 0], 'fp': cm[0, 1],
            'fn': cm[1, 0], 'tp': cm[1, 1]}
#cv_results = cross_validate(model, X_train, y_train, cv=5,scoring= 'f1_macro')
scores = cross_val_score(model, X_train, y_train, cv=5, scoring='f1_macro')
scores
```

Out[134...] array([0.48993069, 0.48750073, 0.48536558, 0.48210634, 0.4843144 ])



In [136...

```

from sklearn.model_selection import ShuffleSplit
n_samples = X.shape[0]
cv = ShuffleSplit(n_splits=5, test_size=0.3, random_state=0)
cross_val_score(model, X, y, cv=cv)

```

Out[136...

```
array([0.88765149, 0.88955126, 0.8887979 , 0.89092696, 0.88683262])
```

In [87]:

```

from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import label_binarize
from sklearn.multiclass import OneVsRestClassifier
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import GridSearchCV
from sklearn.naive_bayes import BernoulliNB
#import sklearn.metrics import accuracy_score
import sklearn.metrics as metrics
nb = BernoulliNB()
param_grid = {'alpha':[1000,100,10,1,0.1,0.01,0.001]} #params we need to try
gsv = GridSearchCV(nb,param_grid,cv=2,verbose=1,n_jobs=-1,scoring='f1')
gsv.fit(X_train,y_train)
nb = BernoulliNB(alpha=0.1)
nb.fit(X_train,y_train)
train_pred = nb.predict(X_train)
cv_pred = nb.predict(X_train)
test_pred = nb.predict(X_test)
y_prob = nb.predict_proba(X_train)
print("Train Set Accuracy: {}".format(metrics.accuracy_score(train_pred, y_tr

```

Fitting 2 folds for each of 7 candidates, totalling 14 fits

Train Set Accuracy: 0.8896344544893031

For our logistic regression model, we have applied three different methods to validate the model: 70/30 split, KNN, and shuffleSplit. 70/30 split method splits the dataframe into 70% train and 30% test. After we predicted using our test features using this cross-validation method, we were able to produce an accuracy rate of 88.79%. Precision is 89%, recall is 100%, f1 is 94% for readmittance, non-readmittance precision 48%, recall 2%, f1 3%

KNN method used to predict on test data using the K using the trained data from nearest Neighbor value (K value). Instead of splitting the data into two parts, data gets split into K parts – 8 parts. This validation process produces accuracy rate of 88.4%. Precision is 89%, recall is 100%, f1 is 94%. For re-admittance, non-readmittance precision 38%, recall 2%, f1 3%

ShuffleSplit method used random samples from the entire dataset, random test and train sets are created during 10 iteration process. The method produces a final accuracy score using the average of 10 accuracy scores from the iteration run. This model produced accuracy score of 88.80%

Given the performance of three of the above methods, we prefer ShuffleSplit method for training our model. This method has produced the top scores for our performance metrics, for example, it has produced the top accuracy score of 88.80%, which is slightly above of 70/30 split method of 88.79%

## Model Interpretability and Explainability

The following are the top ten important features for our model:

Encounter\_id Patient\_nbr gender Discharge\_disposition Admission\_source\_id  
Num\_procedures Num\_medications Number\_outpatient Number\_emergency  
Number\_impatient

These variables are important based on the higher coefficient values of those features. We find these features to be more accurate in predicting the target variable of hospital readmission.

# Conclusion

We notice that cross-validation improves the performance of the model by reducing overfitting. We recommend ShuffleSplit cross-validation method since it has the best performance metrics although we have noticed 70/30 split method's scores are close to that method. In terms of patient demographics, gender is one of the top features out of 10 overall important features of patient re-admittance. Patient visits during both regular and emergency, number of procedures, number of medication consumed by the patients are very important features that hospital may find useful for their operation strategies.

## Recommendation

In addition, we noticed that when recall rate increases, this has an inverse effect on precision and accuracy rate. Therefore, the team recommends that management reduce the recall admittance rate to improve on accuracy, and precision rate. This will save money and resources to cater to new arrivals.

In [ ]: